



SiDi

SAMSUNG

Residência em Engenharia e Ciência de Dados

Processamento de Dados em Larga Escala

{ SPARK }

KDD Cup 99 | PT7 Web

Prof. Jairson B. Rodrigues
jbr@cin.ufpe.br



{ exemplo aplicado / KDDCup 99 }

- Detecção de padrões de ataque em conexões de rede a partir da análise de log conexões disponibilizada pelo MIT Lincoln Labs (KDD Cup 99)
 - <https://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html>
- 4.898.431 de observações
- 41 características (sete categóricas, demais contínuas)
- Classe 0 / normal = 972781, 0,19859%,
- Classe 1 / ataque = 3925650, 0,80141%
- 709.18 Megabytes

{ atributos }

Atributo	Descrição	Tipo
duration	length (number of seconds) of the connection	contínua
protocol_type	type of the protocol, e.g. tcp, udp, etc.	discreta
service	network service on the destination, e.g., http, telnet, etc.	discreta
src_bytes	number of data bytes from source to destination	contínua
dst_bytes	number of data bytes from destination to source	contínua
flag	normal or error status of the connection	discreta
land	1 if connection is from/to the same host/port; 0 otherwise	discreta
wrong_fragment	number of "wrong" fragments	contínua
urgent	number of urgent packets	contínua
hot	number of hot indicators	contínua
num_failed_logins	number of failed login attempts	contínua
logged_in	1 if successfully logged in; 0 otherwise	discreta
num_compromised	number of "compromised" conditions	contínua
root_shell	1 if root shell is obtained; 0 otherwise	discreta
su_attempted	1 if "su root" command attempted; 0 otherwise	discreta
num_root	number of "root" accesses	contínua
num_file_creations	number of file creation operations	contínua
num_shells	number of shell prompts	contínua
num_access_files	number of operations on access control files	contínua
num_outbound_cmds	number of outbound commands in an ftp session	contínua
is_hot_login	1 if the login belongs to the "hot" list; 0 otherwise	discreta
is_guest_login	1 if the login is a "guest" login; 0 otherwise	discreta
count	connections to the same host in the past two seconds	contínua

error_rate	% of connections that have "SYN" errors	contínua
reror_rate	% of connections that have "REJ" errors	contínua
same_srv_rate	% of connections to the same service	contínua
diff_srv_rate	% of connections to different services	contínua
srv_count	connections to the same service in the past two seconds	contínua
srv_error_rate	% of connections that have "SYN" errors	contínua
srv_reror_rate	% of connections that have "REJ" errors	contínua
srv_diff_host_rate	% of connections to different hosts	contínua
label	0 = normal, 1 = ataque	contínua

{ distribuição dos dados }

```
Master Name Node
ubuntu@name-node:~$ wget http://kdd.org/cupfiles/KDDCupData/1999/kddcup.data.zip
--2016-12-11 21:10:35-- http://kdd.org/cupfiles/KDDCupData/1999/kddcup.data.zip
Resolving kdd.org (kdd.org)... 72.10.51.228
Connecting to kdd.org (kdd.org)[72.10.51.228]:80... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: http://www.kdd.org/cupfiles/KDDCupData/1999/kddcup.data.zip [following]
--2016-12-11 21:10:36-- http://www.kdd.org/cupfiles/KDDCupData/1999/kddcup.data.zip
Resolving www.kdd.org (www.kdd.org)... 72.10.51.228
Reusing existing connection to kdd.org:80.
HTTP request sent, awaiting response... 200 OK
Length: 18119640 (17M) [application/zip]
Saving to: 'kddcup.data.zip'

100%[=====>] 18,119,640  10.7MB/s   in 1.6s

2016-12-11 21:10:37 (10.7 MB/s) - 'kddcup.data.zip' saved [18119640/18119640]

ubuntu@name-node:~$ unzip kddcup.data.zip
Archive:  kddcup.data.zip
  inflating: kddcup.data.txt
ubuntu@name-node:~$ ls -l kddcup.data.txt
-rw----- 1 ubuntu ubuntu 742579867 Oct 17  2005 kddcup.data.txt
ubuntu@name-node:~$ hdfs dfs -mkdir /kddcup
ubuntu@name-node:~$ hdfs dfs -mkdir /kddcup/input
ubuntu@name-node:~$ hdfs dfs -put kddcup.data.txt /kddcup/input/
ubuntu@name-node:~$
```

{ simplificação pedagógica :-) }

- Todos os tipos de ataque foram classificados como classe = 1
- Padrões que não correspondem a ataques, classe = 0
- Variáveis categóricas foram submetidas a
 - StringIndexer: converte variáveis categóricas para índices
 - OneHotEncoder: converte índices vetores binários
 - VectorAssembler: agrega todas as características em um único vetor esperso

{ vetores esparsos vs vetores densos }

- Vetor denso:

[1, 0, 66, 29.4, 0, 0, 0, 0, 0, 0, -9, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

- Vetor esparsos:

(26, [0, 2, 3, 10], [1, 66, 29.4, -9])

Tamanho

Índices não nulos

Valores

{ tratamento de dados categóricos }

```
1
2  val indexer: Array[PipelineStage] = originalColumns
3    .map(cname => new StringIndexer()
4      .setInputCol(cname)
5      .setOutputCol(s"${cname}_i")
6      .setHandleInvalid("skip"))
7
8  val one_hot_encoder: Array[PipelineStage] = allIndexedCategoricalColumns
9    .map(cname => new OneHotEncoder()
10      .setInputCol(cname)
11      .setOutputCol(s"${cname}_e"))
12
13  val pipelineTmp = new Pipeline().setStages(indexer ++ one_hot_encoder)
14
15  val df = pipelineTmp.fit(textDF).transform(textDF)
16
17  val assembler = new VectorAssembler()
18    .setInputCols(finalFields.diff(Array("label_i")))
19    .setOutputCol("features")
20
21  val output = assembler.transform(df).select("label_i", "features")
22    .withColumnRenamed("label_i", "label")
23
24  output.write.format("parquet")
25    .save("hdfs://81.14.183.180:9000/kddcup/input/indexed_encoded_data")
```

{ classificador naïve bayes }

```
val random = Math.abs(scala.util.Random.nextInt)
val Array(trainingData, testData) = data.randomSplit(Array(0.7, 0.3), random)
val bayes = new NaiveBayes()
val pipeline = new Pipeline().setStages(Array(bayes))

val paramGrid = new ParamGridBuilder()
  .addGrid(bayes.smoothing, Array(0.0, 0.1, 0.25, 0.5, 0.75, 0.9, 1.0))
  .build()

val cv = new CrossValidator()
  .setEstimator(pipeline)
  .setEvaluator(new BinaryClassificationEvaluator)
  .setEstimatorParamMaps(paramGrid)
  .setNumFolds(10)

val cvModel = cv.fit(trainingData)

val results = cvModel
  .transform(testData)
  .select("prediction", "label")
```


{ classificador naïve bayes }

```
val rows = results.collect()

val numCorrectPredictions = rows
  .map(row => if (row.getDouble(0) == row.getDouble(1)) 1 else 0)
  .foldLeft(0)(_ + _)

val accuracy = 1.0D * numCorrectPredictions / rows.size

val predictionAndLabels = results
  .rdd.map(x => (x(0).asInstanceOf[Double], x(1).asInstanceOf[Double]))

val metrics = new BinaryClassificationMetrics(predictionAndLabels)
```

Acurácia: de 96%,

Área sob a curva Precision-Recall: 0.917

Área sob a curva ROC: 0.902

{ regressão logística }

```
val random = Math.abs(scala.util.Random.nextInt)
val Array(trainingData, testData) = data.randomSplit(Array(0.7, 0.3), random)
val lr = new LogisticRegression().setMaxIter(10)
val pipeline = new Pipeline().setStages(Array(lr))

val paramGrid = new ParamGridBuilder()
  .addGrid(lr.regParam, Array(0.1, 0.01))
  .build()

val cv = new CrossValidator()
  .setEstimator(pipeline)
  .setEvaluator(new BinaryClassificationEvaluator)
  .setEstimatorParamMaps(paramGrid)
  .setNumFolds(10)

val cvModel = cv.fit(trainingData)

val results = cvModel
  .transform(testData)
  .select("prediction", "label")
```

{ regressão logística }

```
val rows = results.collect()

val numCorrectPredictions = rows
  .map(row => if (row.getDouble(0) == row.getDouble(1)) 1 else 0)
  .foldLeft(0)(_ + _)

val accuracy = 1.00 * numCorrectPredictions / rows.size

val predictionAndLabels = results
  .rdd.map(x => (x(0).asInstanceOf[Double], x(1).asInstanceOf[Double]))

val metrics = new BinaryClassificationMetrics(predictionAndLabels)
```

Acurácia: de 97,5%,

Área sob a curva Precision-Recall: 0.947

Área sob a curva ROC: 0.945

{ prática pt7 web corpus }

The screenshot shows the IEEE DataPort website. At the top, there's a navigation bar with 'IEEE DataPort', 'DATASETS', 'COMPETITIONS', 'SUBMIT A DATASET', a search bar, and the IEEE logo. Below this is a blue banner with the word 'Datasets' and 'Open Access' on the right. The main title of the dataset is 'PT7 WEB, AN ANNOTATED PORTUGUESE LANGUAGE CORPUS'. To the left of the title is a blue box containing a list of files: 'pt7-00000.parquet', 'pt7-00001.parquet', 'pt7-00002.parquet', 'pt7-00003.parquet', an ellipsis, 'pt7-00197.parquet', 'pt7-00198.parquet', and 'pt7-00199.parquet'. To the right of the title, there are social media icons for Facebook, Twitter, YouTube, LinkedIn, and Email. Below the title, the 'Citation Author(s)' are listed: 'Jairson Rodrigues (Federal University of Pernambuco)', 'Germano Vasconcelos (Federal University of Pernambuco)', and 'Paulo Madid (Federal University of Pernambuco)'. The 'Submitted by' is 'Jairson Rodrigues'. The 'Last updated' date is 'Sun, 12/06/2020 - 16:48'. The 'DOI' is '10.21227/fhrm-n966'. The 'Data Format' is '.csv (tar.gz)'. The 'License' is 'Creative Commons Attribution 4.0'. There is a '157 Views' icon and text. The 'Categories' are 'Cloud Computing' and 'Machine Learning'. The 'Keywords' are 'big data, Machine Learning algorithms, corpus, spark'. At the bottom left, there's a star rating system with 5 stars and the text '0 ratings - Please login to submit your rating.'. At the bottom right, there are three buttons: 'ACCESS DATASET', 'CITE', and 'SHARE/EMBED'.

IEEE DataPort DATASETS COMPETITIONS SUBMIT A DATASET SEARCH IEEE

Datasets

Open Access

PT7 WEB, AN ANNOTATED PORTUGUESE LANGUAGE CORPUS

pt7-00000.parquet
pt7-00001.parquet
pt7-00002.parquet
pt7-00003.parquet
...
pt7-00197.parquet
pt7-00198.parquet
pt7-00199.parquet

Citation Author(s): Jairson Rodrigues (Federal University of Pernambuco)
Germano Vasconcelos (Federal University of Pernambuco)
Paulo Madid (Federal University of Pernambuco)

Submitted by: Jairson Rodrigues

Last updated: Sun, 12/06/2020 - 16:48

DOI: 10.21227/fhrm-n966

Data Format: .csv (tar.gz)

License: Creative Commons Attribution 4.0

157 Views

Categories: Cloud Computing
Machine Learning

Keywords: big data, Machine Learning algorithms, corpus, spark

0 ratings - Please login to submit your rating.

ACCESS DATASET CITE SHARE/EMBED

<https://dx.doi.org/10.21227/fhrm-n966>

{ aplicações auto-contidas }

- Utilizadas para submeter jobs em spark para o cluster

```
/* SimpleApp.scala */
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
import org.apache.spark.SparkConf

object SimpleApp {
  def main(args: Array[String]) {
    val logFile = "YOUR_SPARK_HOME/README.md" // Should be some file on your system
    val conf = new SparkConf().setAppName("Simple Application")
    val sc = new SparkContext(conf)
    val logData = sc.textFile(logFile, 2).cache()
    val numAs = logData.filter(line => line.contains("a")).count()
    val numBs = logData.filter(line => line.contains("b")).count()
    println(s"Lines with a: $numAs, Lines with b: $numBs")
    sc.stop()
  }
}
```

{ aplicações auto-contidas - compilação }

- Em Scala...
- Estrutura do projeto
 - `./simple.sbt`
 - `./src`
 - `./src/main`
 - `./src/main/scala`
 - `./src/main/scala/SimpleApp.scala`

Compilação

`sbt package`

...

`[info] Packaging {..}/{..}/target/scala-2.11/simple-project_2.11-1.0.jar`

{ aplicações auto-contidas - dependências }

- Arquivo de dependências (simple.sbt)

```
name := "Simple Project"
```

```
version := "1.0"
```

```
scalaVersion := "2.11.7"
```

```
libraryDependencies += "org.apache.spark" %% "spark-core" % "2.1.0"
```

{ aplicações auto-contidas - execução }

- SPARK-SUBMIT (localhost)
YOUR_SPARK_HOME/bin/spark-submit \
--class "SimpleApp" \
--master **local[*]** \
<path>/target/scala-2.11/simple-project_2.11-1.0.jar
- SPARK-SUBMIT (utilizando o cluster em modo YARN)
YOUR_SPARK_HOME/bin/spark-submit \
--class "SimpleApp" \
--master **yarn** \
<path>/target/scala-2.11/kdd-cup_2.11-1.0.jar
- SPARK-SUBMIT (utilizando o cluster em modo Standalone)
YOUR_SPARK_HOME/bin/spark-submit \
--class "SimpleApp" \
--master **spark://spark-node1:7077** \
<path>/target/scala-2.11/kdd-cup_2.11-1.0.jar

{ spark-submit }

- KDDCupETL
- KDDCupRL
- KDDCupNaiveBayes

- **No Namenode**

spark-submit --class "<nome_do_objeto>" \

--master yarn \

--deploy-mode <modo> \

--conf spark.serializer=org.apache.spark.serializer.KryoSerializer \

/home/ubuntu/mlib/target/scala-2.11/kddcup_2.11-1.0.jar

- cluster
- client

{ spark-shell }

```
$ spark-shell
--master spark://spark-node1:7077
--num-executors 2
--driver-memory 6G
--executor-memory 4G
--executor-cores 4
-i KDD-ETL.scala

... OU KDD-NB.scala
... OU KDD-LR.scala
```

{ para onde ir agora? }

- [Machine Learning Library Guide](#)
- [Extração de Características](#)
- [Questões de Otimização](#)
- [Métricas de Avaliação](#)

